

CONSTITUTIONAL OPERATING DOCTRINE

Status

CONSTITUTIONAL_OPERATING_DOCTRINE_V1
UNIFIED_FRAMEWORK_FOR_MACHINE_MEDIATED_PUBLIC_POWER
AI_CLARITY_ACT_EXTENSION
ALMS_MODEL_GOVERNANCE_TRIAGE_BRIDGE

0. Unified Preamble

When government uses automated systems to determine or materially influence liberty, property, benefits, eligibility, enforcement, or access to public opportunity, due process requires more than after-the-fact explanation. The state may not delegate public power to uninspectable processes. If the state acts at machine speed, the affected person must be able to trigger verification at machine speed. Replayable execution proves only what the machine did; it does not prove that the rule was lawful, fair, or constitutional. And when delay functions as denial, queue design itself becomes adjudication. Machine-mediated public power is constitutionally valid only when its decisions are replayable, its governance record is challengeable, and its review path remains meaningful to the person affected.

1. Incorporated Protocols

This doctrine binds five protocols into one operating framework:

1. `alms_citizen_relay_system.md`
2. `MODEL_GOVERNANCE_CHALLENGE.md`
3. `AGENCY_IMPLEMENTATION_CHECKLIST.md`
4. `TRIAGE_AUDIT_PROTOCOL.md`
5. `ai_clarity_act_executive_order_draft.md`

Together they define constitutional due process for machine-mediated state power.

This is not AI ethics.

This is not responsible innovation branding.

This is constitutional procedure adapted to systems that decide, route, score, explain, and enforce at machine speed.

2. Definitions

Term	Definition	Primary Protocol
ALMS	Audit Ledger Manifest Store; the triple-write, content-addressed civic record system for replay manifests and verification results.	ALMS Citizen Relay System
ALMS_DECISION_REPLAY_MANIFEST	Machine-readable record of a covered automated decision, including model identifier, policy hash, transformation chain hash, output receipt hash, and replay requirements.	ALMS Citizen Relay System
APPEND_ONLY_AUDIT	A record structure where new events may be added but prior events may not be silently rewritten or deleted.	ALMS / Agency Checklist
AUTOMATED_ROUTING_AUTHORITY	STATUS SUSPENDED when a triage system loses authority to route cases automatically due to replay or explanation failure.	Triage Audit Protocol
BOTH_TRACKS_REQUIRED	Routing state where a challenge must proceed through both execution fidelity review and model governance review.	Agency Checklist
CITIZEN_REPLAY_RIGHT	The right of an affected person or independent agent to retrieve and inspect replay surfaces without operator permission.	ALMS Citizen Relay System
COVERED_SYSTEM	Any automated or semi-automated system that determines or materially influences public outcomes or protected interests.	AI Clarity Act / ALMS
DELAY_FUNCTIONS_AS_DENIAL	State triggered when queue latency makes review functionally unavailable or meaningless.	Triage Audit Protocol

Term	Definition	Primary Protocol
EXECUTION_FIDELITY_CHALLENGE	Challenge asking whether the system did what the manifest says it did.	ALMS / Implementation Memo
GOVERNANCE_AUTHORITY_SUSPENDED	State entered when a system lacks required governance records or plain-language explainability.	Model Governance Challenge
HIGH_IMPACT_SYSTEM	Covered system affecting benefits, housing, employment, credit, eligibility, legal recommendation, safety-critical access, education, medical access, immigration, or enforcement prioritization.	AI Clarity Act / Model Governance
HUMAN_REVIEW_REQUIRED	State requiring a qualified human official to review the decision, record, or challenge before adverse automated authority continues.	ALMS / Agency Checklist / Triage
LEGACY_REVIEW_REQUIRED	Migration state for pre-ALMS systems that remain active or relied upon but lack full replay surfaces.	ALMS Implementation Memo
MATERIALLY_CONSEQUENTIAL	Capable of affecting liberty, property, benefits, eligibility, housing, credit, employment, legal exposure, public safety access, medical access, education opportunity, or public-sector prioritization.	AI Clarity Act
MODEL_GOVERNANCE_CHALLENGE	Challenge asking whether the system was allowed to do what it did, even if execution replay succeeds.	Model Governance Challenge
NO_OPERATOR_PERMISSION	Verification must not require the system operator's approval, credential, API key, or private dashboard access.	ALMS Citizen Relay System

Term	Definition	Primary Protocol
PATTERN_REVIEW	Review triggered when repeated challenges, explanation failures, replay failures, or upheld complaints indicate systemic risk.	Agency Checklist
PLAIN_LANGUAGE_EXPLANATION	Explanation understandable to the affected person, identifying the rule applied, decisive information, threshold, contestable points, and evidence that could change the result.	Model Governance / Agency Checklist
REPLAY_FAILURE	Failure to reproduce or verify the relevant manifest, output, routing classification, explanation, or audit chain.	ALMS / Triage
REPLAYABLE_EXPLANATION_RECORD	Plain-language explanation tied to a decision ID, manifest hash, policy hash, model version, and explanation hash.	Agency Checklist
STANDARD_HUMAN_QUEUE	Default non-urgent human review queue used when automated triage replay fails or routing authority is suspended.	Triage Audit Protocol
SUSPENDED	State where automated authority pauses pending human review, replay repair, or governance review.	ALMS Citizen Relay System
TRIAL_ROUTING_MANIFEST	ALMS-covered manifest for routing decisions, including triage model ID, policy hash, routing output, confidence, queue ID, and manifest hash.	Triage Audit Protocol
UNGOVERNABLE_EXPLANATION_FAILURE	Entered when a model or triage system cannot explain in plain language why a decision or route occurred.	Model Governance / Triage

Term	Definition	Primary Protocol
UNGOVERNABLE_MODEL_SUNSET	Remedy removing automated authority from a model whose public logic cannot be explained to affected persons.	Model Governance Challenge
VOID_AB_INITIO	Legal consequence for unreplayable punitive or adverse legal outputs: the output is treated as invalid from inception and may not be cited or used downstream.	ALMS Citizen Relay System

3. Citizen Challenge Flow

Readable architecture for the person whose rights are affected:

PERSON SAYS:

"The computer said no and I think that is wrong."

↓

INTAKE ACCEPTS CHALLENGE

Governed by: MODEL_GOVERNANCE_CHALLENGE + AGENCY_IMPLEMENTATION_CHECKLIST

Rule: no technical vocabulary required

↓

ALMS EVIDENCE AUTO-ATTACHED

Governed by: ALMS_CITIZEN_RELAY_SYSTEM

Attached: manifest hash, policy hash, model ID, output receipt, audit log, replay stat

↓

TRIAGE CLASSIFIES ROUTE

Governed by: TRIAGE_AUDIT_PROTOCOL

Rule: triage may classify; triage may not decide

Output: EXECUTION_FIDELITY_CHALLENGE / MODEL_GOVERNANCE_CHALLENGE / BOTH_TRACKS_REQUIRE

↓

TRIAGE ROUTING REPLAYS

Governed by: TRIAGE_AUDIT_PROTOCOL

Citizen can verify: classified as X, version Y, policy hash Z, queue Q
Failure: default to STANDARD_HUMAN_QUEUE or URGENT_HUMAN_REVIEW

↓

REPLAYABLE EXPLANATION GENERATED

Governed by: AGENCY_IMPLEMENTATION_CHECKLIST + MODEL_GOVERNANCE_CHALLENGE
Rule: explanation tied to manifest hash and policy hash
Failure: UNGOVERNABLE_EXPLANATION_FAILURE

↓

TRACK 1: EXECUTION FIDELITY REVIEW

Governed by: ALMS
Question: did the system do what the record says?
Failure: SUSPENDED or VOID_AB_INITIO for punitive/legal outputs

↓

TRACK 2: MODEL GOVERNANCE REVIEW

Governed by: MODEL_GOVERNANCE_CHALLENGE
Question: was the system allowed to do that at all?
Failure: policy revision, threshold review, model withdrawal, or UNGOVERNABLE_MODEL_SU

↓

PUBLIC DASHBOARD UPDATES

Governed by: AGENCY_IMPLEMENTATION_CHECKLIST + TRIAGE_AUDIT_PROTOCOL
Shows: challenge rates, latency, explanation failures, upheld challenges, pattern sign

↓

REMEDY

Possible: human review, corrected decision, suspended authority, void output, policy r

4. Agency Obligations Table

Protocol	Agency Obligation	Responsible Actor	Noncompliance Trigger	Required Remedy
AI Clarity Act	Identify covered and high-impact systems.	Agency head / CIO / program owner	Undisclosed automated system affecting protected interests	Inventory publication and RE-VIEW_REQUIRED
AI Clarity Act	Preserve replay surfaces for materially consequential outputs.	Program owner / system operator	Missing model ID, policy hash, timestamp, input provenance, transformation chain, or output hash	SUSPENDED until manifest is produced
ALMS	Publish decision manifest at time of decision.	System operator	Missing manifest after cutover	SUSPENDED; punitive/legal output VOID_AB_INITIO
ALMS	Triple-write manifest to agency, public records custodian, and content-addressable network.	Records officer / system operator	Missing store copy or hash mismatch	REVIEW_REQUIRED; presumption of invalidity
ALMS	Publish hourly Merkle anchor.	ALMS operator / records custodian	Missing root or timestamp authority	REVIEW_REQUIRED; anchor repair required
ALMS	Permit verification without operator permission.	System operator / vendor	API key, credential, dashboard-only access, or operator gatekeeping required	NO_OPERATOR_PERMISS violation; SUSPENDED
Model Governance	Publish pre-deployment deliberation record.	Agency policy owner	No deliberation record	GOVERNANCE_AUTHORITY

Protocol	Agency Obligation	Responsible Actor	Noncompliance Trigger	Required Remedy
Model Governance	Provide plain-language explanation to affected person.	Program owner / explanation service	Incomprehensible or missing explanation	UNGOVERNABLE_EXPLANATION
Model Governance	Maintain feature, threshold, validation, disparate impact, and drift policy records.	Model governance board / program owner	Missing or defective governance record	MODEL_GOVERNANCE_RECORD
Agency Checklist	Auto-attach ALMS evidence packet at intake.	Challenge intake system	Staff or citizen required to hunt logs	INCOMPLETE_RECORD_REVIEW
Agency Checklist	Ensure triage classification is not disposition.	Intake owner / review supervisor	Auto-denial, auto-dismissal, auto-closure	HUMAN_REVIEW_REQUIRED
Agency Checklist	Publish governance outcomes dashboard.	Agency transparency officer	Challenge patterns hidden or unavailable	PUBLIC_DASHBOARD_UPDATE
Agency Checklist	Trigger pattern review when thresholds are crossed.	Governance board	Repeated failures ignored	PATTERN_REVIEW / UNGOVERNABILITY_REVIEW
Triage Audit	Treat triage as a covered system.	Intake owner / triage operator	Routing model lacks manifest	AUTOMATED_ROUTING_AUDIT
Triage Audit	Make routing outcomes replayable.	Triage operator	Routing cannot be reproduced from version and policy hash	Default to STANDARD_HUMAN_QUEUE or URGENT_HUMAN_REVIEW
Triage Audit	Publish queue latency by routing class.	Agency operations lead	Latency hidden or not broken out by class	DELAY_FUNCTIONS_AS_I
Triage Audit	Explain routing in plain language.	Triage operator	Cannot explain why routed to queue	UNGOVERNABLE_EXPLANATION suspend automated routing

5. Two-Track Doctrine

Track 1: Execution Fidelity

Execution fidelity asks:

Did the system do what the manifest says it did?

It is governed by ALMS.

Evidence includes:

- decision replay manifest
- manifest hash
- output receipt hash
- model identifier
- policy hash
- transformation chain hash
- triple-write persistence
- Merkle anchor
- audit log
- replay result

Execution failure may trigger:

SUSPENDED

REVIEW_REQUIRED

VOID_AB_INITIO

HUMAN_REVIEW_REQUIRED

Track 2: Substantive Legitimacy

Substantive legitimacy asks:

Was the system allowed to do that at all?

It is governed by MODEL_GOVERNANCE_CHALLENGE.

Evidence includes:

- deliberation record
- legal authority
- plain-language logic
- feature governance
- threshold policy
- drift tolerance policy
- validation record
- disparate impact analysis
- public comment record
- human review pathway

Governance failure may trigger:

MODEL_GOVERNANCE_REVIEW
POLICY_REVISION_REQUIRED
THRESHOLD_REVIEW_REQUIRED
FEATURE_REVIEW_REQUIRED
GOVERNANCE_AUTHORITY_SUSPENDED
UNGOVERNABLE_MODEL_SUNSET
PUBLIC_RULEMAKING_REQUIRED
COURT_REVIEW_AVAILABLE

The Firewall

REPLAY PROVES EXECUTION.
IT DOES NOT PROVE JUSTICE.

A valid hash does not cure an invalid rule.

A perfect replay does not cure an unlawful threshold.

A stable model does not cure an ungovernable explanation.

A timely queue does not cure a routing classifier that cannot be replayed.

The Reunification at Remedy

The tracks separate diagnosis but reunify at remedy.

A single challenge may require:

- execution correction
- human review
- governance review
- public dashboard update
- pattern review
- court review
- suspension of automated authority

The citizen does not need to choose the correct legal or technical track.

The institution must route the challenge correctly and preserve review.

6. Closing Doctrine

NO MANIFEST, NO AUTHORITY.
NO REPLAY, NO AUTOMATED POWER.
NO DELIBERATION RECORD, NO AUTOMATED AUTHORITY.
NO PLAIN-LANGUAGE EXPLANATION, NO AUTOMATED AUTHORITY.
TRIAGE MAY CLASSIFY; TRIAGE MAY NOT DECIDE.
A DELAY THAT FUNCTIONS AS A DENIAL IS A DECISION.

A SCORE IS NOT A VERDICT.
CONSENSUS IS NOT CONVERGENCE.

The technology may change.

The rights do not.

This doctrine exists so public power implemented by machines remains public, reviewable, challengeable, and bounded by constitutional due process. # AGENCY OPPOSITION BRIEF

Strongest Government Opposition to Replayability-Based Procedural Due Process Claim

Draft template for agency opposition to motion for summary judgment or motion to dismiss.

This model brief is a stress-test document. It presents the strongest non-caricatured government arguments against converting replayable manifests, speed parity, and ALMS-style infrastructure into constitutional requirements.

IN THE UNITED STATES DISTRICT COURT

Jane Doe, Plaintiff,

v.

United States Department of Health and Human Services, Defendant.

Civil Action No. _____

DEFENDANT’S MEMORANDUM IN OPPOSITION TO PLAINTIFF’S MOTION FOR SUMMARY JUDGMENT AND IN SUPPORT OF DISMISSAL

INTRODUCTION

Plaintiff asks this Court to impose a sweeping new constitutional infrastructure mandate on federal benefits administration. She frames the request as ordinary procedural due process, but the relief sought would require agencies to build and preserve replayable machine-level manifests for automated or semi-automated decisions, restructure legacy systems, expose decision logic to public replay, and shift the burden of proof in routine benefits adjudications.

The Constitution does not require that. The Administrative Procedure Act does not require that. Congress has not required that.

The agency provided Plaintiff with notice, an opportunity to appeal, and access to the administrative record that exists. Plaintiff's theory would transform the absence of a preferred technical artifact into a constitutional violation. But courts review agency action on the record before the agency; they do not order agencies to redesign adjudication infrastructure because a different record architecture might be desirable.

Automated tools do not eliminate due process. But due process has never required agencies to preserve every intermediate computational step in a publicly replayable form. The requested rule would disrupt benefits administration nationwide, impose billions of dollars in modernization burdens, create fraud and security risks, and transfer complex policy choices from Congress and agencies to the judiciary.

The motion should be denied, and the complaint should be dismissed or remanded under existing administrative procedures, not constitutionalized into an ALMS mandate.

ARGUMENT

I. The APA and Due Process Do Not Require Agencies to Create Records That Do Not Exist.

Plaintiff's central claim is that the agency failed to provide a "replayable manifest." But no statute, regulation, or constitutional rule required the agency to create such a manifest at the time of decision. The administrative record consists of the materials actually before the agency when it acted. The APA does not require agencies to reconstruct records that were never generated or to design systems that would generate Plaintiff's preferred technical artifacts.

Citizens to Preserve Overton Park v. Volpe, 401 U.S. 402 (1971), does not help Plaintiff. Overton Park requires judicial review based on the administrative record; it does not authorize courts to dictate the architecture of that record in advance. The agency produced the denial notice, benefits history, eligibility records, and case materials relied upon in the decision. Plaintiff may contest those materials through ordinary administrative appeal.

Plaintiff attempts to convert a record-format objection into a due-process violation. But the Constitution has never required the government to create a perfect evidentiary record before taking administrative action. Due process requires notice and an opportunity to be heard appropriate to the context. It does not require cryptographic manifests, hash chains, model identifiers, or machine-replay infrastructure.

If the agency's explanation is inadequate, ordinary remand may be appropriate. But that is a far narrower remedy than Plaintiff seeks. The Court should not hold that the absence of a non-required technical manifest creates a presumption of constitutional invalidity.

II. Mathews Requires Balancing, Not Speed-Parity Absolutism.

Plaintiff invokes *Mathews v. Eldridge*, 424 U.S. 319 (1976), but her proposed rule collapses the balancing test into a categorical technical mandate. Mathews requires courts to weigh the private interest, the risk of erroneous deprivation and value of additional safeguards, and the government’s interest, including fiscal and administrative burdens.

The government’s interest here is substantial. Federal benefits systems process millions of eligibility determinations, renewals, changes, appeals, and data matches. Many systems are decades old, federally integrated with state systems, and dependent on statutory timelines. Retrofitting each materially consequential action with replayable manifests, triple-write storage, public anchoring, and independent replay access would require years of procurement, system redesign, privacy analysis, security review, staffing, and congressional appropriations.

Plaintiff calls this “speed parity.” In practice, it is an infrastructure command that would slow or halt benefits processing for everyone. Every additional precondition on automated processing carries a cost in delayed enrollment, delayed renewals, and delayed payment. Mathews does not permit courts to ignore those burdens by declaring that machine-speed decisions must always be machine-replayable.

The Constitution permits reasonable administrative systems, not only technically ideal ones. Existing procedures provide notice, appeal rights, human review, and judicial review. The fact that Plaintiff would prefer a more granular record does not make current procedures unconstitutional.

III. Public Replayability Creates Serious Security, Privacy, and Program-Integrity Risks.

Plaintiff’s proposed manifest regime would require agencies to expose decision logic, thresholds, input categories, model versions, transformation chains, and replay paths. That information may be useful to auditors, but it is also useful to adversaries.

Benefits systems must prevent fraud, identity manipulation, false eligibility claims, data poisoning, and coordinated attacks on public programs. Revealing how an automated system weights inputs, applies thresholds, detects anomalies, or routes cases can enable gaming. Even if source code is withheld, detailed replay manifests may disclose enough operational logic to compromise program integrity.

The risks are not limited to fraud. Per-decision manifests may reveal sensitive personal data, protected medical or disability information, household composition, immigration-sensitive facts, income data, or law-enforcement-adjacent indicators. Public or broadly accessible replay surfaces create attack surfaces. Requiring verification without operator permission may conflict with privacy law, cybersecurity obligations, anti-fraud duties, and statutory confidentiality provisions.

The agency can support controlled audit, inspector general review, internal quality assurance, and privacy-protected appeals. But Plaintiff’s demand for citizen-triggered, operator-permission-free replay goes far beyond what due process requires. Aggregated audit, confi-

dential review, and human appeal can protect both fairness and security without exposing operational details to the public.

IV. Plaintiff's Theory Raises Separation-of-Powers and Major-Questions Concerns.

Plaintiff asks the Court to establish a nationwide rule: no materially consequential automated decision without a replayable manifest. That rule would restructure federal adjudication infrastructure across benefits, housing, credit, eligibility, enforcement, immigration, medical access, and safety systems.

Such a mandate is a legislative policy choice. Congress knows how to impose recordkeeping, audit, technology, cybersecurity, privacy, and benefits-administration requirements. It has not enacted ALMS. It has not required triple-write manifests, public content-addressable anchoring, independent replay access, or machine-speed verification. Courts should not impose those requirements as constitutional law.

The major questions doctrine cautions against discovering broad regulatory authority in ambiguous statutory or constitutional language when the issue has vast economic and political significance. Plaintiff's proposed rule would require massive appropriations, procurement, cybersecurity redesign, personnel retraining, and intergovernmental coordination. These are choices for Congress and agencies, not a district court applying due process in an individual benefits case.

Due process doctrine is flexible, but flexibility does not authorize the judiciary to become a national systems architect. The Court should reject Plaintiff's attempt to constitutionalize an elaborate technical framework under the label of meaningful review.

V. The Agency Provided Constitutionally Adequate Process Under Existing Law.

Plaintiff received notice of the denial and the opportunity to appeal. She may submit evidence, challenge factual assumptions, request human review, and seek judicial review after exhausting administrative remedies. Those procedures satisfy due process in the benefits context, particularly given the government's need to administer programs efficiently and consistently.

Plaintiff's disagreement with the level of technical detail in the notice does not establish a constitutional violation. Due process does not require the agency to disclose every internal analytical step, especially where disclosure would burden administration or compromise program integrity.

If the Court concludes that the notice should have been more specific, the proper remedy is limited remand for clarification or further explanation, not a presumption that all automated decisions lacking replay manifests are unconstitutional.

VI. The Requested Burden Shift Is Unsupported.

Plaintiff asks the Court to shift the burden to the agency to prove the automated decision was correct whenever a replayable manifest is absent. That rule would invert ordinary administrative review. Agencies are entitled to a presumption of regularity. A plaintiff challenging agency action bears the burden of showing legal error.

The absence of a novel technical artifact does not justify stripping the agency of that presumption. Otherwise, any litigant could demand a preferred form of record architecture and claim procedural violation when it is missing. That would destabilize administrative law and incentivize collateral attacks on system design rather than focused review of the actual benefits determination.

VII. Plaintiff's Proposed Two-Track Doctrine Is Policy Argument, Not Binding Law.

Plaintiff distinguishes “execution fidelity” from “substantive legitimacy.” That may be a useful academic framework. But it is not a constitutional test. Courts already review agency action for legal error, factual support, reasoned decisionmaking, and procedural adequacy. The APA and due-process doctrine do not require courts to adopt Plaintiff's terminology or infrastructure.

The Court should not mistake a policy architecture for a legal rule. The question is whether Plaintiff received constitutionally adequate process under existing law. She did.

CONCLUSION

The government does not contend that automated adjudication is immune from review. It contends that due process does not require the judiciary to impose a comprehensive replayability infrastructure that Congress has not enacted and agencies have not been funded or directed to build.

The agency provided notice, appeal rights, and an administrative record. If additional explanation is needed, ordinary remand is available. But Plaintiff's proposed rebuttable presumption, speed-parity doctrine, public replay mandate, and ALMS-style infrastructure requirement exceed the role of the Court and the requirements of the Constitution.

The motion for summary judgment should be denied. The complaint should be dismissed or, at most, remanded for further explanation under ordinary administrative law principles.

Respectfully submitted,

Counsel for Defendant

Appendix A: Government’s Proposed Rule

An agency’s failure to create a replayable technical manifest for an automated or semi-automated benefits decision does not itself establish a procedural due-process violation absent a statute, regulation, or specific constitutional requirement mandating such a manifest.

Courts may require adequate explanation and reviewable records under existing law, but should not impose prospective technical infrastructure mandates as constitutional rules.

Appendix B: Limited Alternative Relief

If the Court finds the agency’s explanation inadequate, the appropriate remedy is limited remand for further explanation or human review under existing administrative procedures, without adopting a rebuttable presumption of invalidity for all automated decisions lacking replayable manifests. # REPLY BRIEF

Plaintiff’s Reply in Support of Procedural Due Process for Automated Benefits Denial

Draft template for reply in support of motion for summary judgment or petition for review of agency action.

This model reply responds point-by-point to the strongest agency opposition to replayability, speed parity, and manifest-based review. It frames the doctrine as existing constitutional procedure applied to machine-mediated adjudication, not as AI ethics or a new technology policy mandate.

IN THE UNITED STATES DISTRICT COURT

Jane Doe, Plaintiff,

v.

United States Department of Health and Human Services, Defendant.

Civil Action No. _____

PLAINTIFF’S REPLY IN SUPPORT OF MOTION FOR SUMMARY JUDGMENT

INTRODUCTION

The agency’s opposition rests on a single premise: because it chose not to preserve the computational record of an automated benefits denial, the Court must review only the paper fragments the agency elected to keep. That premise is wrong.

When an agency delegates a protected-interest decision to a computational process, that process is the decisional path. The administrative record must contain enough of that path to permit meaningful review. A denial letter is not the record of a machine-mediated decision. It is a summary of the result. If the agency cannot show how the result was produced, it has not produced the record; it has produced an assertion.

Plaintiff does not ask the Court to enact artificial-intelligence policy. Plaintiff asks the Court to apply existing due process principles to the actual facts of modern adjudication. Notice, hearing, and judicial review cannot be meaningful when the operative decision process is uninspectable. The Constitution does not require Congress to speak again before courts enforce the Fifth Amendment against a new administrative mechanism.

The motion should be granted.

ARGUMENT

I. The Agency’s “Record That Exists” Argument Is a Category Error.

The agency argues that courts review the record that exists, not the record Plaintiff wishes existed. But this framing mistakes an omission for a record.

When a human adjudicator makes a decision, the record includes the materials, reasoning, and path sufficient to permit review. When a machine makes or materially shapes the decision, the computational process is the decisional path. Omitting that path from the record is not a neutral record-compilation choice. It is a failure to preserve how the decision was made.

Overton Park requires review of the whole administrative record. The whole record of a machine-mediated decision includes the machine’s logic as executed, at least to the extent necessary to identify the rule or model version, policy version, inputs or input categories, transformation chain, output integrity record, and explanation tied to the actual decision. A PDF denial letter is not that record. It is a conclusion.

The agency says no manifest existed. That is the problem, not the answer. An agency cannot avoid review by designing a decisional process that leaves no reviewable trace. The absence of a record of the operative process is not a feature of the system entitled to deference. It is the procedural defect.

The Court need not require perfection. But it must require reviewability. A record that cannot reveal what rule was applied, what version governed, what information mattered,

or how the output was produced is not a meaningful administrative record for a protected-interest deprivation.

II. This Case Does Not Present a Major Question; It Presents Ordinary Due Process.

The agency attempts to recast Plaintiff’s claim as a nationwide infrastructure mandate. That is not the relief requested and not the rule the Court needs to adopt.

Procedural due process is not a major question. It is the constitutional floor. *Mathews* and *Goldberg* do not lie dormant until Congress instructs agencies how to preserve records for automated adjudication. Courts have always adapted procedural guarantees to the mechanism of deprivation before them.

Plaintiff does not ask the Court to order ALMS, triple-write storage, content-addressable networks, or any specific architecture. ALMS is an example of a standard capable of satisfying due process in machine-mediated administration. The legal rule is simpler: when an agency chooses to automate a materially consequential decision, it must preserve a record sufficient to inspect that automated decision.

How the agency satisfies that requirement is for the agency. Whether the Constitution permits an uninspectable protected-interest deprivation is for the Court.

The agency’s separation-of-powers argument therefore fails. The Court is not being asked to become a systems architect. It is being asked to do what courts do: decide whether the process used to deprive Plaintiff of a protected interest was constitutionally adequate.

III. The Agency Weaponizes Efficiency Against Mathews.

The agency argues that replayability would burden administration and reduce automation’s efficiency. That argument gets the third *Mathews* factor backwards.

Administrative efficiency is a legitimate interest. But efficiency gained by eliminating reviewability is not a constitutional interest. It is the mechanism of the due process violation. The government cannot reduce process, increase opacity, act faster, and then cite the speed it gained from opacity as the reason courts should tolerate the risk.

The second *Mathews* factor—the risk of erroneous deprivation—is not static. It scales with speed, opacity, and population. A single human error may affect one file. A single model error, policy-version error, data-ingestion defect, or drift event can affect thousands before any human notices. In that environment, the risk is structural, not marginal.

When risk becomes structural, additional safeguards have greater value, not less. A replayable decisional record directly reduces the risk of erroneous deprivation by allowing the agency, citizen, hearing officer, and court to determine what happened. Without it, review becomes guesswork.

Nor are the agency’s claimed efficiency savings self-proving. A black-box system that produces unreviewable denials may appear efficient until it generates appeals, remands, injunc-

tions, class actions, emergency reviews, and loss of public trust. One replayable system that survives review is cheaper than an opaque system that cannot explain itself.

Speed parity is not a demand that every citizen personally rerun federal software. It is the principle that when the state acts at machine speed, the review path cannot remain trapped in paper-speed opacity. The affected person must be able to trigger verification of the decisional record within a meaningful time.

IV. Security Concerns Do Not Justify Unaccountability.

The agency conflates public disclosure with verifiability.

Plaintiff does not demand publication of model weights, fraud thresholds, source code, anomaly-detection rules, or sensitive personal data to the world. Plaintiff demands a reviewable record sufficient for the affected person, counsel, agency reviewer, and court to understand and test the basis of the decision.

Courts routinely manage sensitive records. They handle classified materials, trade secrets, sealed medical records, juvenile records, law-enforcement materials, and confidential business information. Protective orders, redactions, in camera review, secure expert access, summaries, and controlled disclosure are ordinary tools of judicial administration. The existence of sensitive information does not eliminate due process.

Security through obscurity is not a constitutional defense. If the agency cannot defend its process without hiding the operative basis from the person harmed by it, the process is not secure; it is unaccountable. Those are different properties.

A replayable manifest need not expose fraud thresholds to adversaries. It can identify the policy version, decision path, input categories, output integrity, and explanation necessary for review while protecting specific sensitive details. The agency’s argument proves only that implementation must be careful. It does not prove that reviewability may be denied.

V. Reasonableness Is Contextual, and Uninspectable Machine-Speed Deprivation Is Not Reasonable.

The agency says the Constitution permits reasonable systems, not technically ideal ones. Plaintiff agrees. The question is what reasonableness requires when the state chooses machine-mediated adjudication.

In 1970, a paper file and an in-person hearing could be reasonable. In 2026, when a protected-interest decision is generated or materially shaped by an automated system, a paper summary that omits the operative computational path is not reasonable. It is arbitrary by design.

Goldberg guarantees a meaningful opportunity to be heard. A person cannot meaningfully contest a decision whose basis cannot be examined. *Mathews* requires attention to the risk of erroneous deprivation and the value of safeguards. When decisions are fast, opaque, and repeatable at population scale, the safeguard of a replayable decisional record is not technical perfection; it is the minimum condition for meaningful review.

The agency chose machine speed. Having chosen it, the agency cannot claim that preserving a record of what the machine did is unreasonable. The Constitution does not forbid automation. It forbids the government from using automation to place protected-interest decisions beyond inspection.

VI. The Proper Rule Is Narrow and Practical.

The Court need not decide the future of federal automation. It need only decide this case.

The appropriate rule is:

When an agency delegates a materially consequential decision affecting a protected liberty or property interest to an automated or semi-automated system, procedural due process requires a reviewable record sufficient to determine what the system did and why the affected person was harmed.

Where the agency cannot produce such a record, the decision is presumptively procedurally defective, and the agency bears the burden of producing a meaningful explanation or providing prompt human review.

That rule does not ban automation. It does not constitutionalize a specific vendor architecture. It does not require public disclosure of sensitive fraud rules. It simply prevents the government from using a machine to make a decision and then telling the citizen, the hearing officer, and the court to trust the output.

VII. The Two-Track Doctrine Is Existing Review Logic, Not New Law.

The agency calls Plaintiff's execution-fidelity and substantive-legitimacy distinction academic. It is not. It is the ordinary difference between asking whether the agency did what it says it did and whether what it did was lawful.

Execution fidelity asks whether the system executed as recorded. Substantive legitimacy asks whether the rule, threshold, feature use, policy, or model design was lawful and constitutionally sufficient.

Courts already distinguish factual record, procedural regularity, reasoned decisionmaking, and legal validity. The two-track doctrine simply gives that familiar structure a vocabulary suited to machine-mediated decisions.

A replayable record proves execution. It does not prove justice. A valid hash does not cure an invalid rule. But without the record of execution, the Court cannot even reach meaningful review of the rule.

VIII. Delay That Functions as Denial Is Reviewable Due Process Injury.

The agency suggests ordinary appeal rights are enough. But a right to appeal into an indefinite queue is not meaningful process.

Where routing or triage determines whether a person receives review in days or years, delay itself becomes materially consequential. A delay that functions as denial is a decision. A decision requires due process.

The agency cannot deny benefits at machine speed and then offer review at backlog speed without producing the record needed to review the denial. If the agency lacks that record, remand must be prompt, human, and meaningful.

PROPOSED HOLDING

This Court should hold that when an agency delegates materially consequential decisions to automated or semi-automated systems, due process requires:

1. a replayable decisional record establishing what the system did;
2. the ability of the affected person to contest either the fidelity of execution or the lawfulness of the rule; and
3. recognition that delay functionally equivalent to denial triggers the same protections as denial itself.

These are not new rights. They are the existing procedural guarantees of the Fifth Amendment applied to the facts of machine-speed governance.

CONCLUSION

The agency asks the Court to treat the absence of a computational record as a reason to defer. It is the reason not to defer.

The government may use machines. It may not use machines to make protected-interest decisions beyond meaningful review. A denial that cannot be replayed cannot be tested. A process that cannot be inspected cannot satisfy due process. And a queue that delays review until the harm is complete is not review at all.

Plaintiff's motion should be granted.

NO MANIFEST, NO AUTHORITY.

NO REPLAY, NO AUTOMATED POWER.

NO PLAIN-LANGUAGE EXPLANATION, NO AUTOMATED AUTHORITY.

A DELAY THAT FUNCTIONS AS A DENIAL IS A DECISION.

Respectfully submitted,

Counsel for Plaintiff